

▼ PHSAE 4 PROJECT - TEAM 3

Introduction

This is the team 3 phase 3 project notebook. Our group members include :

- Anthony Ngatia
- Jessyca Aperi
- Joy Chepchumba
- Naomi Rotich

In this notebook we shall be analysing the **Zillow Housing Dataset using a Time Series Model**

Overview

In these project we shall seek to do the following :

- Load the dataset
- Understand the dataset
- Choose our target variable.
- Prepare the dataset (Example : Cleaning the dataset , checking for multicollinearity)
- Encode our categorical variables
- Make several models
- Evaluate our models
- Use our models for prediction
- Come up with relevant findings.

▼ 1). Business Understanding

Real estate investment is a lucrative and dynamic industry that requires careful analysis and decision-making. The fictional real estate investment firm is seeking guidance on identifying the top 5 zip codes for investment opportunities. To address this question, historical data from Zillow Research is utilized.

i) Background:

Real estate investment is a lucrative and dynamic industry that requires careful analysis and decision-making. The fictional real estate investment firm is seeking guidance on identifying the top 5 zip codes for investment opportunities. To address this question, historical data from Zillow Research is utilized. The dataset contains information on various attributes, including RegionID, RegionName, City, State, Metro, SizeRank, CountyName, and value (real estate prices).

ii). Main Objective:

The main objective of this project is to identify the top 5 zip codes that offer the best investment potential in terms of real estate prices. By analyzing historical trends and patterns, the project aims to provide actionable insights to the investment firm, enabling them to make informed decisions on where to allocate their resources.

Specific Objectives:

- **Analyze Historical Data:** The project involves analyzing the historical data of real estate prices across different zip codes. This includes understanding the trends, patterns, and fluctuations in property values over time.
- **Identify Promising Zip Codes:** Using the analysis of historical data, the project aims to identify the zip codes that have shown consistent growth, stability, or potential for future appreciation. These zip codes are considered the most favorable for investment.
- **Consider Location Factors:** In addition to the historical performance, the project also takes into account location-specific factors such as city, state, and metro. This information helps assess the overall desirability and attractiveness of the investment opportunities.
- **Evaluate Market SizeRank:** The SizeRank attribute provides insights into the relative size and competitiveness of the real estate market in each zip code. This factor helps gauge the potential opportunities and risks associated with investing in a particular area.

2). Data Understanding

The dataset contains information on various attributes, including RegionID, RegionName, City, State, Metro, SizeRank, CountyName, and value (real estate prices). Our dataset is the Zillow Housing Dataset which was sourced from Zillow Research Page.

In order to understand how our dataset looks like lets get a preview of this data by loading it.

Importing the Zillow Housing Dataset

```
#Importing data libraries
import numpy as np
import pandas as pd

#importing visualisation libraries
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

#importing math libraries
from math import sqrt

#Importing modeling libraries
from statsmodels.tsa.seasonal import seasonal_decompose
from dateutil.parser import parse
import statsmodels.api as sm
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
from matplotlib.pylab import rcParams
from statsmodels.tsa.arima.model import ARIMA
from sklearn.metrics import mean_squared_error as MSE

#importing the dataset
df = pd.read_csv('/content/zillow_data.csv')
df.head()
```

	RegionID	RegionName	City	State	Metro	CountyName	SizeRank	1996-04	1996-05	1996-06	...	2017-07	2017-08	2017-09	2017-10
0	84654	60657	Chicago	IL	Chicago	Cook	1	334200.0	335400.0	336500.0	...	1005500	1007500	1007800	1008000
1	90668	75070	McKinney	TX	Dallas-Fort Worth	Collin	2	235700.0	236900.0	236700.0	...	308000	310000	312500	314000
2	91982	77494	Katy	TX	Houston	Harris	3	210400.0	212200.0	212200.0	...	321000	320600	320200	320000
3	84616	60614	Chicago	IL	Chicago	Cook	4	498100.0	500900.0	503100.0	...	1289800	1287700	1287400	1287000
4	93144	79936	El Paso	TX	El Paso	El Paso	5	77300.0	77300.0	77300.0	...	119100	119400	120000	120000

5 rows x 272 columns



```
#Investigating for the shape of the dataset
df.shape

(14723, 272)
```

The dataset contains 14723 rows and 272 columns

```
#Describing the dataset
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14723 entries, 0 to 14722
Columns: 272 entries, RegionID to 2018-04
dtypes: float64(219), int64(49), object(4)
memory usage: 30.6+ MB
```

The dataset contains data types of; float, integers and strings

3). Data Preparation

Checking for missing values

```
df.isna().sum()

RegionID      0
RegionName    0
City          0
State         0
Metro        1043
...
2017-12       0
2018-01       0
2018-02       0
2018-03       0
2018-04       0
Length: 272, dtype: int64
```

#Displaying the rows with missing values.

```
df[df.isnull().any(axis=1)].iloc[:1]
```

	RegionID	RegionName	City	State	Metro	CountyName	SizeRank	1996-04	1996-05	1996-06	...	2017-07	2017-08	2017-09
20	61625	10011	New York	NY	New York	New York	21	NaN	NaN	NaN	...	12137600	12112600	120
36	61796	10456	New York	NY	New York	Bronx	37	NaN	NaN	NaN	...	357900	357100	3
105	84613	60611	Chicago	IL	Chicago	Cook	106	NaN	NaN	NaN	...	1475200	1473900	14
151	69340	27410	Greensboro	NC	NaN	Guilford	152	137100.0	136600.0	136000.0	...	212900	213200	2
156	62048	11238	New York	NY	New York	Kings	157	NaN	NaN	NaN	...	2673300	2696700	27
...
14706	59046	3215	Waterville Valley	NH	Claremont	Grafton	14707	NaN	NaN	NaN	...	786000	780900	7
14707	69681	28039	East Spencer	NC	Charlotte	Rowan	14708	NaN	NaN	NaN	...	27300	26400	1
14708	99401	97733	Crescent	OR	Klamath Falls	Klamath	14709	NaN	NaN	NaN	...	197700	203700	2
14710	59210	3812	Bartlett	NH	NaN	Carroll	14711	80900.0	80800.0	80800.0	...	215500	217000	2
14717	62697	12720	Bethel	NY	NaN	Sullivan	14718	62500.0	62600.0	62700.0	...	122200	122700	1

1827 rows × 272 columns

Dealing with Missing Values

We used Linear Interpolation to handle the numerical missing values

```
# Checking for missing values in each column
missing_columns = df.columns[df.isnull().any()].tolist()

print(missing_columns)

['Metro', '1996-04', '1996-05', '1996-06', '1996-07', '1996-08', '1996-09', '1996-10', '1996-11', '1996-12', '1997-01', '1997-02', '1997-03', '1997-04', '1997-05', '1997-06', '1997-07', '1997-08', '1997-09', '1997-10', '1997-11', '1997-12']

# Iterate over the columns with missing values
for col in missing_columns:
    # Perform interpolation using linear method
    df[col] = df[col].interpolate(method='linear')

# Print the updated dataset with interpolated values
```

```
print(df)
```

	RegionID	RegionName	City	State	Metro	\
0	84654	60657	Chicago	IL	Chicago	
1	90668	75070	McKinney	TX	Dallas-Fort Worth	
2	91982	77494	Katy	TX	Houston	
3	84616	60614	Chicago	IL	Chicago	
4	93144	79936	El Paso	TX	El Paso	
...	
14718	58333	1338	Ashfield	MA	Greenfield Town	
14719	59107	3293	Woodstock	NH	Claremont	
14720	75672	40404	Berea	KY	Richmond	
14721	93733	81225	Mount Crested Butte	CO	NaN	
14722	95851	89155	Mesquite	NV	Las Vegas	

	CountyName	SizeRank	1996-04	1996-05	1996-06	...	2017-07	\
0	Cook	1	334200.0	335400.0	336500.0	...	1005500	
1	Collin	2	235700.0	236900.0	236700.0	...	308000	
2	Harris	3	210400.0	212200.0	212200.0	...	321000	
3	Cook	4	498100.0	500900.0	503100.0	...	1289800	
4	El Paso	5	77300.0	77300.0	77300.0	...	119100	
...	
14718	Franklin	14719	94600.0	94300.0	94000.0	...	216800	
14719	Grafton	14720	92700.0	92500.0	92400.0	...	202100	
14720	Madison	14721	57100.0	57300.0	57500.0	...	121800	
14721	Gunnison	14722	191100.0	192400.0	193700.0	...	662800	
14722	Clark	14723	176400.0	176300.0	176100.0	...	333800	

	2017-08	2017-09	2017-10	2017-11	2017-12	2018-01	2018-02	2018-03	\
0	1007500	1007800	1009600	1013300	1018700	1024400	1030700	1033800	
1	310000	312500	314100	315000	316600	318100	319600	321100	
2	320600	320200	320400	320800	321200	321200	323000	326900	
3	1287700	1287400	1291500	1296600	1299000	1302700	1306400	1308500	
4	119400	120000	120300	120300	120300	120300	120500	121000	
...	
14718	217700	218600	218500	218100	216400	213100	209800	209200	
14719	208400	212200	215200	214300	213100	213700	218300	222700	
14720	122800	124600	126700	128800	130600	131700	132500	133000	
14721	671200	682400	695600	695500	694700	706400	705300	681500	
14722	336400	339700	343800	346800	348900	350400	353000	356000	

	2018-04
0	1030600
1	321800
2	329900
3	1307000
4	121500
...	...
14718	209300
14719	225800
14720	133400
14721	664400
14722	357200

[14723 rows x 272 columns]

```
df = df.drop('Metro', axis = 1)
```

```
#Confirming the missing values are not present
df[df.isnull().any(axis=1)].iloc[: -1]
```

RegionID	RegionName	City	State	CountyName	SizeRank	1996-04	1996-05	1996-06	1996-07	...	2017-07	2017-08	2017-09	2017-10	2017-11	2017-12	2018-01
0 rows x 271 columns																	

Reshaping our dataset from wide to long format

```
#Converting our dataframe time column from float to datetime format

def get_datetimes(df):
    """
    Takes a dataframe:
    returns only those column names that can be converted into datetime objects
    as datetime objects.
    NOTE number of returned columns may not match total number of columns in passed dataframe
    """
```

```

"""

return pd.to_datetime(df_new.columns.values[1:], format='%Y-%m')

# Reshaping our dataset from Wide to Long Format

def melt_data(data):
    melted = pd.melt(data, id_vars=['RegionID', 'RegionName', 'City', 'State', 'SizeRank', 'CountyName'], var_name='time')
    melted['time'] = pd.to_datetime(melted['time'], infer_datetime_format=True)
    melted = melted.dropna(subset=['value'])
    return melted.groupby(['RegionID', 'RegionName', 'City', 'State', 'SizeRank', 'CountyName', 'time']).aggregate({'value': 'mean'}).reset_i

#Printing the first five rows of the long format dataset.
df1 = melt_data(df)
df1.head()

```

	RegionID	RegionName	City	State	SizeRank	CountyName	time	value
0	58196	1001	Agawam	MA	5851	Hampden	1996-04-01	113100.0
1	58196	1001	Agawam	MA	5851	Hampden	1996-05-01	112800.0
2	58196	1001	Agawam	MA	5851	Hampden	1996-06-01	112600.0
3	58196	1001	Agawam	MA	5851	Hampden	1996-07-01	112300.0
4	58196	1001	Agawam	MA	5851	Hampden	1996-08-01	112100.0

```

#Descriptive Statistics for the Long format dataset
df1.describe()

```

	RegionID	RegionName	SizeRank	value
count	3.901595e+06	3.901595e+06	3.901595e+06	3.901595e+06
mean	8.107501e+04	4.822235e+04	7.362000e+03	2.060636e+05
std	3.193304e+04	2.935833e+04	4.250165e+03	2.368017e+05
min	5.819600e+04	1.001000e+03	1.000000e+00	1.130000e+04
25%	6.717400e+04	2.210100e+04	3.681000e+03	9.770000e+04
50%	7.800700e+04	4.610600e+04	7.362000e+03	1.469000e+05
75%	9.092100e+04	7.520600e+04	1.104300e+04	2.354000e+05
max	7.538440e+05	9.990100e+04	1.472300e+04	1.931490e+07

```

# Checking for missing values in the long format dataframe
df1.isna().sum()

```

```

RegionID      0
RegionName    0
City          0
State         0
SizeRank      0
CountyName    0
time          0
value        0
dtype: int64

```

```

#checking for duplicates in the Long Format Dataframe
df1.duplicated().sum()

```

```
0
```

4). Exploratory Data Analysis

Grouping the data by months

```
#Setting the time column as the index
df1.set_index('time', inplace = True)
```


```
#Resampling using monthly buckets
monthlyvalue = df1.resample('MS')
month_mean = monthlyvalue.mean()
month_mean.tail()
```

	RegionID	RegionName	SizeRank	value
time				
2017-12-01	81075.010052	48222.348706	7362.0	281095.320247
2018-01-01	81075.010052	48222.348706	7362.0	282657.060382
2018-02-01	81075.010052	48222.348706	7362.0	284368.688447
2018-03-01	81075.010052	48222.348706	7362.0	286511.376757
2018-04-01	81075.010052	48222.348706	7362.0	288039.944305

Visualising the Dataframe

Grouping per Month and plotting

```
# Plotting the monthly housing average
plt.figure()
plt.plot(month_mean.index, month_mean['value'])
plt.title('Line plot of the Monthly Average house values since 1996')
plt.xlabel('Date')
plt.ylabel('Mean Value')
plt.grid(True)
plt.show()
```




Grouping per Year and plotting

```
#Setting the year as the index
df1['Year'] = df1.index.year
```


```
#Plotting the Housing annual average
df1_yearly = df1['value'].resample('A').mean()
df1_yearly.plot();
plt.title('Line plot of the Annual Average house values since 1996')
plt.xlabel('year')
```

```
plt.ylabel('value')
plt.grid()
```




Grouping per Quarter and plotting

```
#Plotting the Quarterly annual housing average
df1_quarterly = df1['value'].resample('Q').mean()
df1_quarterly.plot();
plt.title('Line plot of the Quarterly Annual Average house values since 1996')
plt.xlabel('Quarter')
plt.ylabel('Value')
plt.show()
```



Grouping per Decade and plotting

```
#Plotting the Average housing values per decade
df1_decade = df1['value'].resample('10Y').mean()
df1_decade.plot();
plt.title('Average housing values per decade since 1996')
plt.xlabel('decade')
plt.ylabel('value')
plt.show()
```



We proceeded to group our dataframe with the Region Name and Value columns only inorder to determine the Region Names with the highest mean_values.

```
# Group the data by the 'RegionName' and 'Year' columns and calculate the mean value:
grouped_data = df1.groupby(['RegionName'])['value'].mean()
grouped_data.head()
```

```
RegionName
1001    174509.811321
1002    273152.452830
1005    172650.943396
1007    217938.113208
1008    175319.622642
Name: value, dtype: float64
```

```
#Resetting the index
grouped_new_df = grouped_data.reset_index()
grouped_new_df.columns = ['RegionName', 'mean_value']
grouped_new_df.head()
```

	RegionName	mean_value
0	1001	174509.811321
1	1002	273152.452830
2	1005	172650.943396
3	1007	217938.113208
4	1008	175319.622642

```
#Sorting the grouped dataframe using the mean_value, in descending order
grouped_new = grouped_new_df.sort_values(by = 'mean_value', ascending = False)
grouped_new.head()
```

	RegionName	mean_value
1405	10021	8.438275e+06
1403	10011	5.444482e+06
1406	10128	5.085436e+06
1404	10014	4.507875e+06
13590	94027	3.487129e+06

```
# slicing the top 10 values
top_10_df = grouped_new.head(10)
top_10_df
```

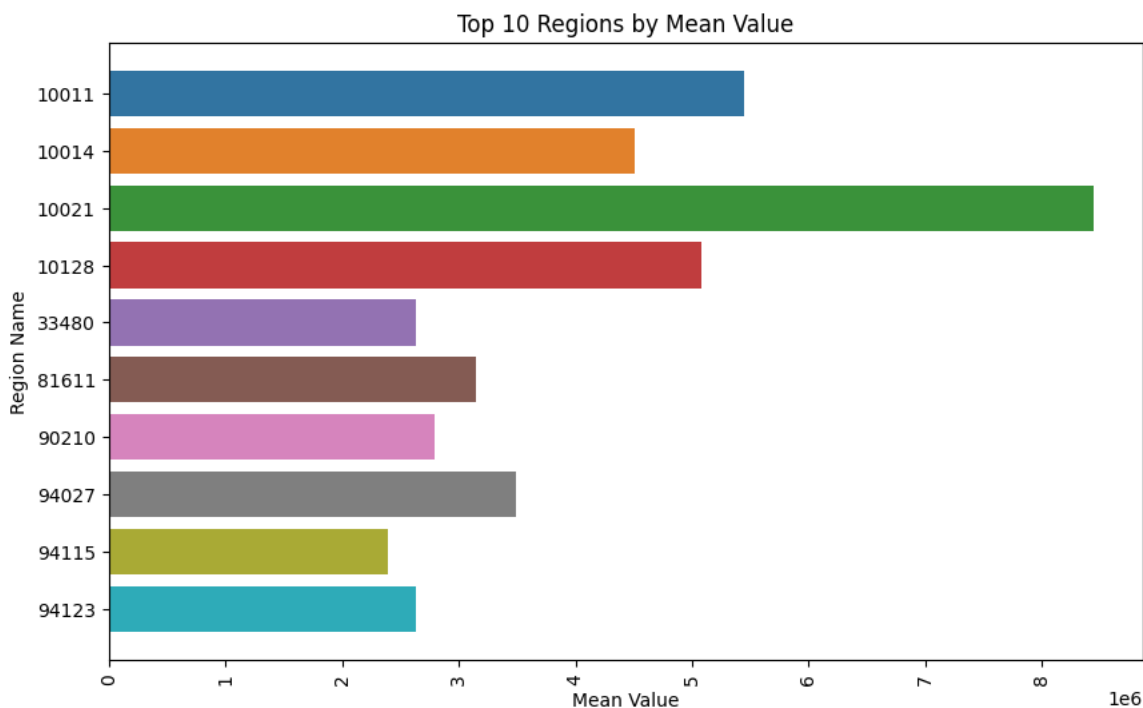

	RegionName	mean_value
1405	10021	8.438275e+06
1403	10011	5.444482e+06
1406	10128	5.085436e+06
1404	10014	4.507875e+06
13590	94027	3.487129e+06
12180	81611	3.147124e+06
12902	90210	2.789977e+06
5528	33480	2.634498e+06
13621	94123	2.630977e+06
13615	94115	2.399030e+06

The output above contains a list of the top 10 Region Names according to the mean_value

We proceeded to visualise this result below with a horizontal bar graph.

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(10, 6))
sns.barplot(x= 'mean_value', y='RegionName', data=top_10_df, orient='h')
plt.xlabel('Mean Value')
plt.ylabel('Region Name')
plt.title('Top 10 Regions by Mean Value')
plt.xticks(rotation=90)
plt.autoscale(enable=True)
plt.show()
```



We chose our metric of investment as Return of Investment (ROI). Defined using the formula below:

$$ROI = \text{Net Investment Gain} / \text{Cost of Investment} \times 100$$

Considering the investment firm is new in the market, we opted to model our dataset using the last 3 years.

```
# Subsetting the data to include only the years 2015 to 2018
start_date = '2015-04-01'
end_date = '2018-04-01'
subset_df = df1.loc[start_date:end_date]
```

```
# Print the subsetting DataFrame
print(subset_df)
```

time	RegionID	RegionName	City	State	SizeRank	CountyName	\
2015-04-01	58196	1001	Agawam	MA	5851	Hampden	
2015-05-01	58196	1001	Agawam	MA	5851	Hampden	
2015-06-01	58196	1001	Agawam	MA	5851	Hampden	
2015-07-01	58196	1001	Agawam	MA	5851	Hampden	
2015-08-01	58196	1001	Agawam	MA	5851	Hampden	
...
2017-12-01	753844	29486	Summerville	SC	3188	Dorchester	
2018-01-01	753844	29486	Summerville	SC	3188	Dorchester	
2018-02-01	753844	29486	Summerville	SC	3188	Dorchester	
2018-03-01	753844	29486	Summerville	SC	3188	Dorchester	
2018-04-01	753844	29486	Summerville	SC	3188	Dorchester	

time	value	Year
2015-04-01	192200.0	2015
2015-05-01	192400.0	2015
2015-06-01	192100.0	2015
2015-07-01	191500.0	2015
2015-08-01	191000.0	2015
...
2017-12-01	182700.0	2017
2018-01-01	183300.0	2018
2018-02-01	184400.0	2018
2018-03-01	186500.0	2018
2018-04-01	188300.0	2018

[544751 rows x 8 columns]

```
# Grouping the melted dataframe by 'RegionName' and 'Year' and calculating the mean
grouped_df = pd.DataFrame(subset_df.groupby(['RegionName', 'Year'])['value'].mean())
grouped_df.columns = ['mean_value']
grouped_df.head(10)
```

RegionName	Year	mean_value
1001	2015	192322.222222
	2016	199033.333333
	2017	212866.666667
	2018	222425.000000
1002	2015	316555.555556
	2016	316950.000000
	2017	333133.333333
	2018	348950.000000
1005	2015	176522.222222
	2016	191550.000000

Grouping the dataset and Calculating the ROI

```
# Group the melted dataframe by 'RegionName' and calculate the mean value for the entire timeframe
grouped_df = subset_df.groupby('RegionName')['value'].mean().reset_index()
grouped_df.columns = ['RegionName', 'mean_value']
```

```
# Calculate ROI by taking percent change of the mean 'value' column
grouped_df['ROI3'] = grouped_df['mean_value'].pct_change(periods=3) * 100
```

```
# Drop the first row since it will have NaN value for ROI
grouped_df = grouped_df.dropna()
```

```
# Sort the DataFrame by ROI in descending order
```

```
grouped_df = grouped_df.sort_values('ROI3', ascending=False)
```

```
# Print the resulting DataFrame with ROI value
print(grouped_df.head(10))
```

	RegionName	mean_value	ROI3
1405	10021	1.844019e+07	5782.191013
1403	10011	1.162878e+07	5578.042150
5038	31561	2.391924e+06	3637.381757
1404	10014	9.444808e+06	2915.375522
12180	81611	4.106662e+06	1667.189644
5703	34102	2.665270e+06	1255.234588
1774	11975	2.934378e+06	1234.414906
12181	81615	2.622397e+06	1019.635130
12809	89413	2.030743e+06	956.295953
14060	96141	6.153486e+05	949.453791


We selected the top 10 RegionNames based on the ROI value as listed below;

- 10021
- 10011
- 31561
- 10014
- 81611
- 34102
- 11975
- 81615
- 89413
- 96141

```
# Select the top 10 rows
top_10_df = grouped_df.head(10)
```

```
# plotting Region names with the highest ROI
# Plot the bar chart
plt.figure(figsize=(10, 6))
sns.barplot(x=top_10_df['RegionName'], y= top_10_df['ROI3'])
plt.xlabel('Region Name')
plt.ylabel('ROI3')
plt.title('Bar chart of the Top 10 Region Names by ROI')
plt.xticks(rotation=90)
plt.autoscale(enable=True)
plt.show()
```

Bar chart of the Top 10 Region Names by ROI



Displaying the location of the top 10 Region Names

```
#Get Location Names
best10_zipcodes = list(top_10_df.RegionName.values)
for i in best10_zipcodes:
    city = df1[df1['RegionName']==i].City.values[0]
    state = df1[df1['RegionName']==i].State.values[0]
    print(f'Zipcode : {i} \nLocation: {city}, {state}\n')

    Zipcode : 10021
    Location: New York, NY

    Zipcode : 10011
    Location: New York, NY

    Zipcode : 31561
    Location: Sea Island, GA

    Zipcode : 10014
    Location: New York, NY

    Zipcode : 81611
    Location: Aspen, CO

    Zipcode : 34102
    Location: Naples, FL

    Zipcode : 11975
    Location: Wainscott, NY

    Zipcode : 81615
    Location: Snowmass Village, CO

    Zipcode : 89413
    Location: Glenbrook, NV

    Zipcode : 96141
    Location: Homewood, CA

#Creating a list of the top 10 Region Names
region_name_list = top_10_df['RegionName'].unique().tolist()
region_name_list

[10021, 10011, 31561, 10014, 81611, 34102, 11975, 81615, 89413, 96141]

# Create a list of region names you want to filter
region_names =[10021, 10011, 31561, 10014, 81611,34102, 11975, 81615, 89413, 96141]

# Filter the original DataFrame based on the region names
filtered_df = subset_df[subset_df['RegionName'].isin(region_names)]

# Create a new DataFrame with only the 'RegionName' and 'value' columns
new_df = filtered_df.loc[:, ['RegionName', 'value']]

# Print the new DataFrame
print(new_df.tail(10))
print(len(new_df))
print(new_df['RegionName'].unique())

      RegionName    value
time
2017-07-01      96141  671300.0
2017-08-01      96141  671500.0
2017-09-01      96141  666500.0
2017-10-01      96141  667100.0
2017-11-01      96141  671800.0
2017-12-01      96141  675000.0
2018-01-01      96141  675000.0
2018-02-01      96141  677500.0
2018-03-01      96141  684400.0
2018-04-01      96141  689700.0
370
[10011 10014 10021 11975 31561 34102 81611 81615 89413 96141]
```

```
#Creating a new df for the yearly sampled values and setting the index as time
new_df['Year'] = new_df.index.year
new_df.head()
```

	RegionName	value	Year
time			
2015-04-01	10011	10572500.0	2015
2015-05-01	10011	10569500.0	2015
2015-06-01	10011	10674900.0	2015
2015-07-01	10011	10848100.0	2015
2015-08-01	10011	11131200.0	2015

```
# Drop the year column
new_df.drop('Year',axis=1,inplace=True)
new_df.head()
```

	RegionName	value
time		
2015-04-01	10011	10572500.0
2015-05-01	10011	10569500.0
2015-06-01	10011	10674900.0
2015-07-01	10011	10848100.0
2015-08-01	10011	11131200.0

Calculating the Descriptive Statistics for the top 10 Region Names in the yearly sampled dataset

```
for region_name1 in new_df['RegionName'].unique():
    region_data = new_df[new_df['RegionName'] == region_name1]

    print(f'Value descriptive statistics for region name {region_name1}:')
    print(region_data['value'].describe())
    print()
```

```
Value descriptive statistics for region name 10011:
count    3.700000e+01
mean     1.162878e+07
std      4.438828e+05
min      1.056950e+07
25%      1.137880e+07
50%      1.173390e+07
75%      1.199880e+07
max      1.213760e+07
Name: value, dtype: float64
```

```
Value descriptive statistics for region name 10014:
count    3.700000e+01
mean     9.444808e+06
std      3.304039e+05
min      8.876100e+06
25%      9.135300e+06
50%      9.458600e+06
75%      9.701500e+06
max      9.958800e+06
Name: value, dtype: float64
```

```
Value descriptive statistics for region name 10021:
count    3.700000e+01
mean     1.844019e+07
std      6.335105e+05
min      1.664400e+07
25%      1.830710e+07
50%      1.852730e+07
75%      1.885970e+07
max      1.931490e+07
Name: value, dtype: float64
```

```
Value descriptive statistics for region name 11975:
count    3.700000e+01
```

```

mean    2.934378e+06
std     3.292249e+05
min     2.343300e+06
25%     2.692000e+06
50%     2.938700e+06
75%     3.207900e+06
max     3.473300e+06
Name: value, dtype: float64

```

Value descriptive statistics for region name 31561:

```

count    3.700000e+01
mean     2.391924e+06
std      8.126152e+04
min      2.262200e+06
25%      2.338800e+06
50%      2.378800e+06
75%      2.453900e+06
max      2.542700e+06
Name: value, dtype: float64

```

Value descriptive statistics for region name 34102:

```

count    3.700000e+01
mean     2.391924e+06
std      8.126152e+04
min      2.262200e+06
25%      2.338800e+06
50%      2.378800e+06
75%      2.453900e+06
max      2.542700e+06
Name: value, dtype: float64

```


▼ Checking for trend

```

#checking for trends in the dataset
for region_name, region_data in new_df.groupby('RegionName'):
    values = region_data['value'].values.flatten() # Flatten the values to create a 1-dimensional array
    plt.plot(values, label=region_name) # Plot the values with region name as label

plt.xlabel('Time')
plt.ylabel('Value')
plt.title('Line Plot by Region Name')
plt.legend()
plt.show()

```



Visualising using line plots for each of the Region Names in the top 10

```

# Group the dataframe by 'RegionName'
grouped_df1 = new_df.groupby('RegionName')

# Create a dictionary to store the separate dataframes for each region name
region_dataframes = {}

# Iterate over each group and create a separate dataframe for each region name
for region_name, region_group in grouped_df1:
    region_dataframes[region_name] = region_group.copy()

# Access the separate dataframes for each region name

```

```
for region_name, region_df in region_dataframes.items():
    print(f'Dataframe for Region Name: {region_name}')
    print(region_df.head())
    print()
```

```
Dataframe for Region Name: 10011
      time  RegionName  value
2015-04-01    10011  10572500.0
2015-05-01    10011  10569500.0
2015-06-01    10011  10674900.0
2015-07-01    10011  10848100.0
2015-08-01    10011  11131200.0
```

```
Dataframe for Region Name: 10014
      time  RegionName  value
2015-04-01    10014  9938600.0
2015-05-01    10014  9827500.0
2015-06-01    10014  9571200.0
2015-07-01    10014  9278700.0
2015-08-01    10014  9252000.0
```

```
Dataframe for Region Name: 10021
      time  RegionName  value
2015-04-01    10021  16644000.0
2015-05-01    10021  16659500.0
2015-06-01    10021  17149200.0
2015-07-01    10021  17775200.0
2015-08-01    10021  17965800.0
```

```
Dataframe for Region Name: 11975
      time  RegionName  value
2015-04-01    11975  2343300.0
2015-05-01    11975  2371200.0
2015-06-01    11975  2394400.0
2015-07-01    11975  2413700.0
2015-08-01    11975  2437600.0
```

```
Dataframe for Region Name: 31561
      time  RegionName  value
2015-04-01    31561  2453900.0
2015-05-01    31561  2482800.0
2015-06-01    31561  2511200.0
2015-07-01    31561  2530500.0
2015-08-01    31561  2531200.0
```

```
Dataframe for Region Name: 34102
      time  RegionName  value
2015-04-01    34102  2481500.0
2015-05-01    34102  2502200.0
2015-06-01    34102  2522100.0
2015-07-01    34102  2529700.0
2015-08-01    34102  2541600.0
```

```
Dataframe for Region Name: 81611
      time  RegionName  value
2015-04-01    81611  3956500.0
```

Since we want to plot the line plots of time against value, we proceeded to drop the RegionNames

```
# Drop the column in each dataframe
# Iterate over each dataframe in the dictionary and drop the 'RegionName' column
for region_name, region_df in region_dataframes.items():
    region_df.drop('RegionName', axis=1, inplace=True)

# Access the modified dataframes
for region_name, region_df in region_dataframes.items():
    print(f'Dataframe for Region Name: {region_name}')
    print(region_df.head())
    print()
```

```
Dataframe for Region Name: 10011
      time  value
2015-04-01    10572500.0
```

```

2015-04-01 10572500.0
2015-05-01 10569500.0
2015-06-01 10674900.0
2015-07-01 10848100.0
2015-08-01 11131200.0

```

```

Dataframe for Region Name: 10014
      value

```

```

time
2015-04-01 9938600.0
2015-05-01 9827500.0
2015-06-01 9571200.0
2015-07-01 9278700.0
2015-08-01 9252000.0

```

```

Dataframe for Region Name: 10021
      value

```

```

time
2015-04-01 16644000.0
2015-05-01 16659500.0
2015-06-01 17149200.0
2015-07-01 17775200.0
2015-08-01 17965800.0

```

```

Dataframe for Region Name: 11975
      value

```

```

time
2015-04-01 2343300.0
2015-05-01 2371200.0
2015-06-01 2394400.0
2015-07-01 2413700.0
2015-08-01 2437600.0

```

```

Dataframe for Region Name: 31561
      value

```

```

time
2015-04-01 2453900.0
2015-05-01 2482800.0
2015-06-01 2511200.0
2015-07-01 2530500.0
2015-08-01 2531200.0

```

```

Dataframe for Region Name: 34102
      value

```

```

time
2015-04-01 2481500.0
2015-05-01 2502200.0
2015-06-01 2522100.0
2015-07-01 2529700.0
2015-08-01 2541600.0

```

```

Dataframe for Region Name: 81611
      value

```

```

time
2015-04-01 3956500.0

```

```



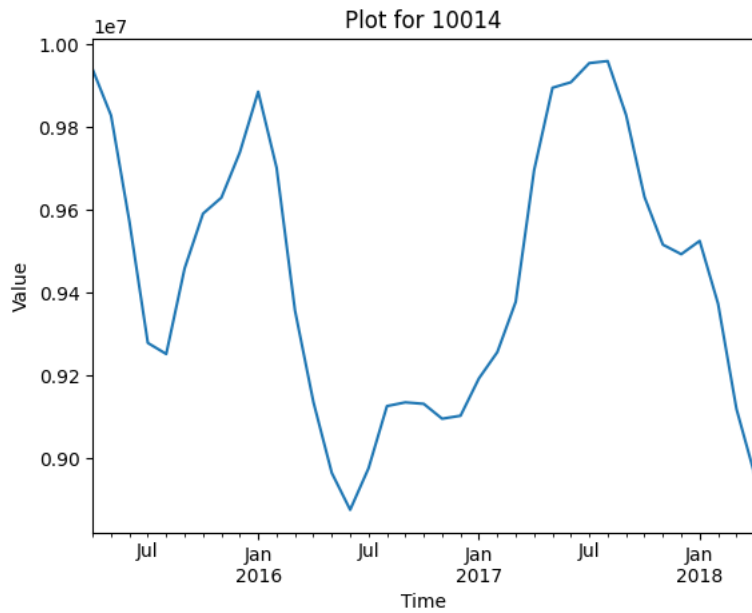
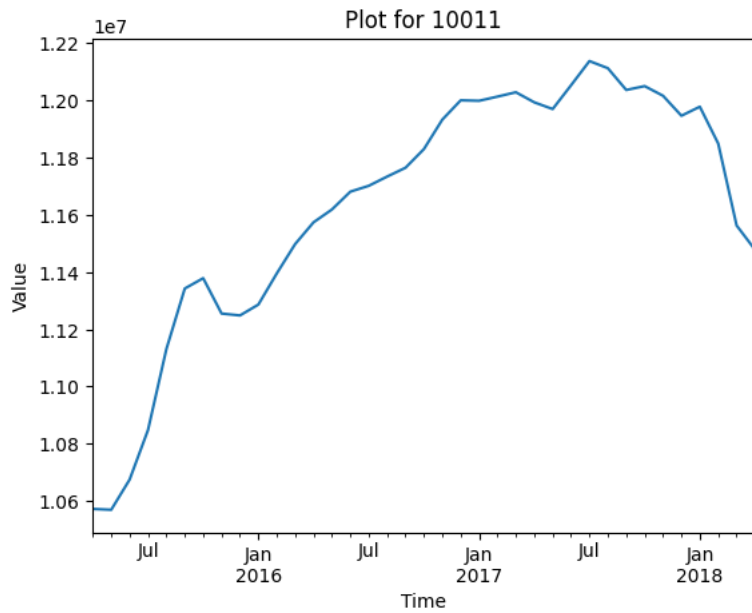
#Iterating over the list of RegionNames and plotting for each


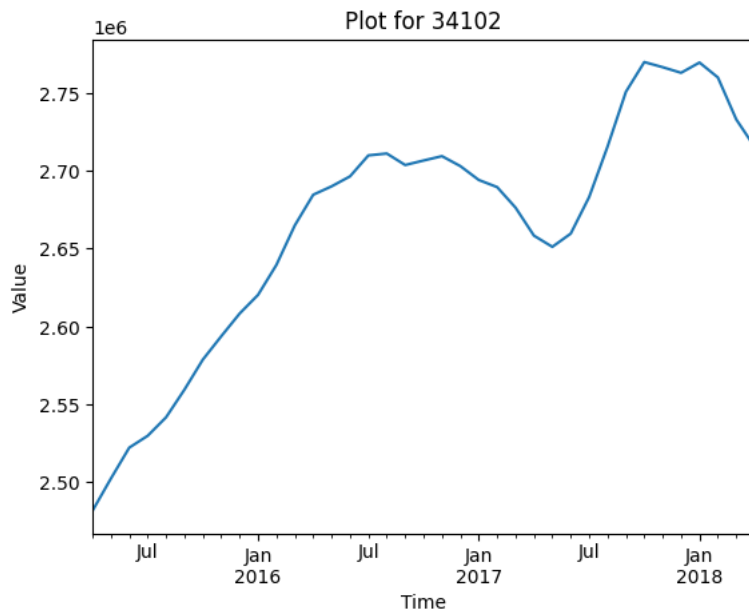
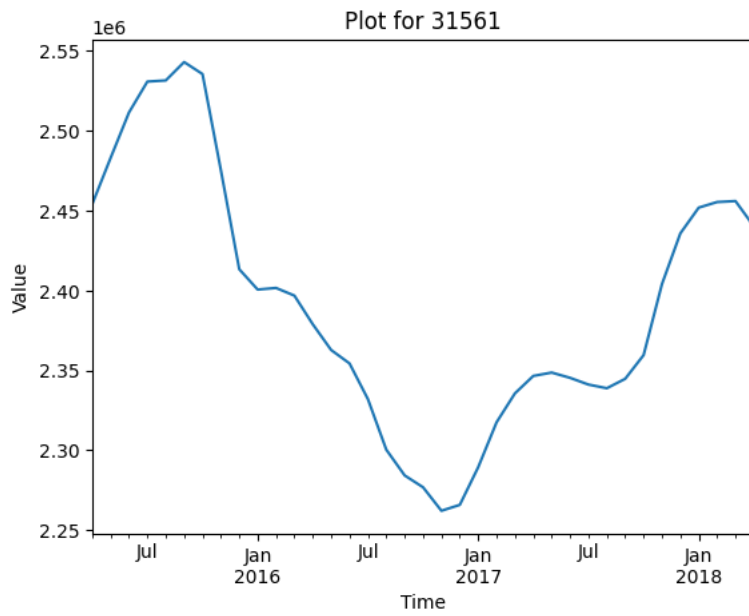

```


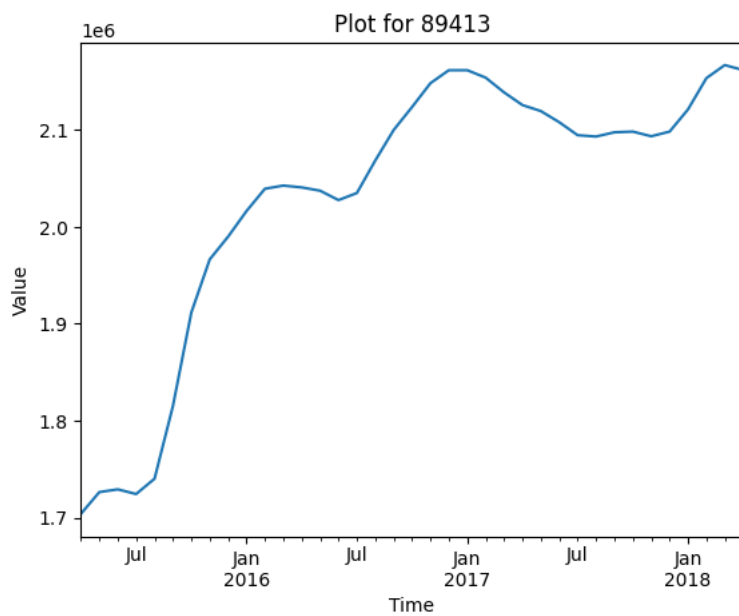
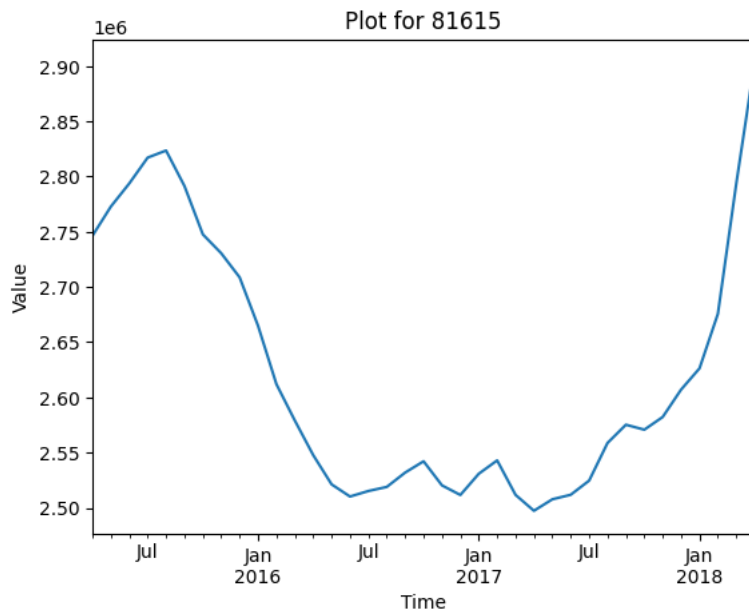

```

for key, df in region_dataframes.items():
    # Plot the 'value' column for the current dataframe
    df['value'].plot()
    plt.title(f'Plot for {key}')
    plt.xlabel('Time')
    plt.ylabel('Value')
    plt.show()

```





640000 ↓

/

|

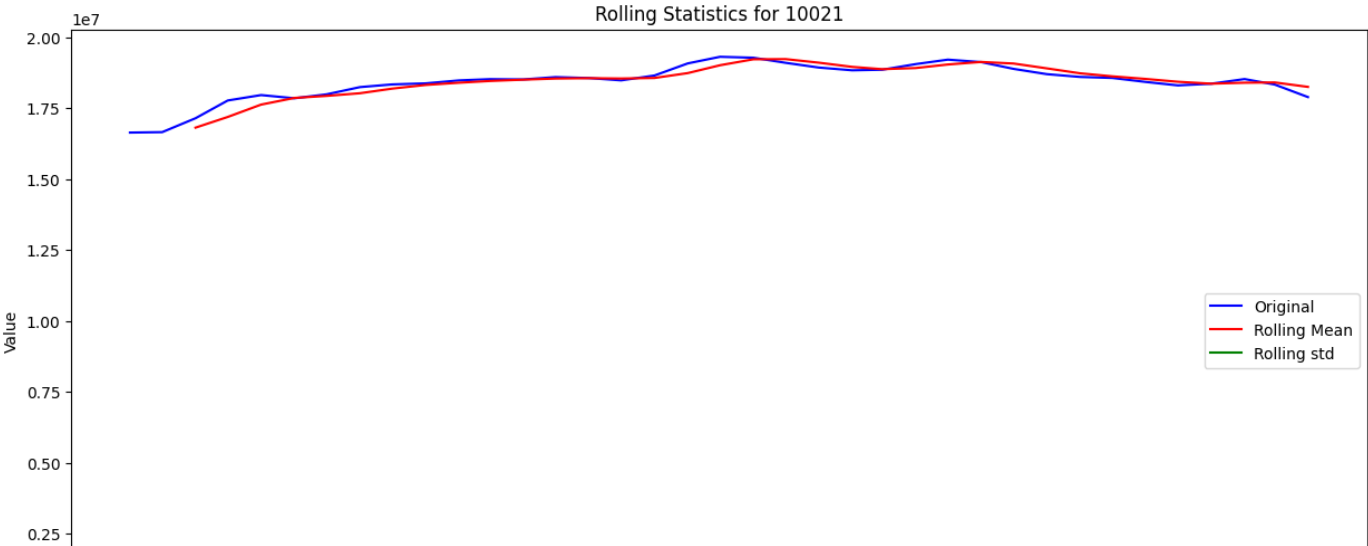


▼ Performing Rolling Statistics





← | /





```
# Define the window size for rolling calculations
window_size = 3


for key, df in region_dataframes.items():
    # Calculate the rolling mean and standard deviation for the 'value' column
    rolling_mean = df['value'].rolling(window=window_size).mean()
    rolling_std = df['value'].rolling(window=window_size).std()

    # Plot the original values, rolling mean, and rolling standard deviation
    fig = plt.figure(figsize=(15, 7))
    plt.plot(df['value'], c='blue', label='Original')
    plt.plot(rolling_mean, c='red', label='Rolling Mean')
    plt.plot(rolling_std, c='green', label='Rolling std')
    plt.xlabel('Months')
    plt.ylabel('Value')
    plt.title(f"Rolling Statistics for {key}")
    plt.legend()
    plt.show()
```









Rolling Statistics for 96141



Performing Dickey Fuller Test to verify the plots above

```

600000 |
#Dickey-Fuller test to verify your visual result.
from statsmodels.tsa.stattools import adfuller

for key, df in region_dataframes.items():
    # Perform Dickey-Fuller test
    print(f'Results of Dickey-Fuller Test for {key}:')
    dfctest = adfuller(df['value'])

    # Extract and display test results
    dfctest = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value', '#Lags Used', 'Number of Observations Used'])
    for k, v in dfctest[4].items():
        dfctest[f'Critical Value ({k})'] = v
    print(dfctest)
    print('\n')

Results of Dickey-Fuller Test for 10011:
Test Statistic      0.837760
p-value             0.992222
#Lags Used          10.000000
Number of Observations Used  26.000000
Critical Value (1%)  -3.711212
Critical Value (5%)  -2.981247
Critical Value (10%) -2.630095
dtype: float64

Results of Dickey-Fuller Test for 10014:
Test Statistic      -3.831754
p-value             0.002599
#Lags Used           8.000000
Number of Observations Used  28.000000
Critical Value (1%)  -3.688926
Critical Value (5%)  -2.971989
Critical Value (10%) -2.625296
dtype: float64

Results of Dickey-Fuller Test for 10021:
Test Statistic      -0.322143
p-value             0.922309
#Lags Used           8.000000
Number of Observations Used  28.000000
Critical Value (1%)  -3.688926
Critical Value (5%)  -2.971989
Critical Value (10%) -2.625296
dtype: float64

Results of Dickey-Fuller Test for 11975:
Test Statistic      -1.169849
p-value             0.686430
#Lags Used           2.000000
Number of Observations Used  34.000000
Critical Value (1%)  -3.639224
Critical Value (5%)  -2.951230
Critical Value (10%) -2.614447
dtype: float64

```



```
Results of Dickey-Fuller Test for 31561:
Test Statistic      -2.571365
p-value             0.099061
#Lags Used          1.000000
Number of Observations Used 35.000000
Critical Value (1%)  -3.632743
Critical Value (5%)  -2.948510
Critical Value (10%) -2.613017
dtype: float64
```

```
Results of Dickey-Fuller Test for 34102:
Test Statistic      -2.339657
p-value             0.159496
```

Interpreting the results:

- 10011 - The test statistic (0.837760) is greater than all the critical values, therefore we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 10014 - The test statistic (-3.831754) is lower than the critical values, and the p-value (0.002599) is less than 0.05, therefore we reject the null hypothesis of non-stationarity. The data provides sufficient evidence to suggest that the time series is stationary.
- 10021 - The test statistic (-0.322143) is greater than the critical values, and the p-value (0.922309) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 11975 - The test statistic (-1.169849) is greater than the critical values, and the p-value (0.686430) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 31561 - Since the test statistic (-2.571365) is greater than the critical values, and the p-value (0.099061) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 34102 - Since the test statistic (-2.339657) is greater than the critical values, and the p-value (0.159496) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 81611 - Since the test statistic (2.577411) is greater than the critical values, and the p-value (0.999071) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 81615 - Since the test statistic (-0.671517) is greater than the critical values, and the p-value (0.853990) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 89416 - Since the test statistic (-2.852473) is greater than the critical values, and the p-value (0.051160) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.
- 96141 - Since the test statistic (-2.464834) is greater than the critical values, and the p-value (0.124267) is greater than 0.05, we fail to reject the null hypothesis of non-stationarity. The data does not provide sufficient evidence to suggest that the time series is stationary.

Based on the above interpretation, all other Zipnames contain non-stationarity apart from 10014. We proceed to perform detrending.

Checking for Stationarity

```
# Defining the check_stationarity function
def stationarity_check(TS):

    # Import adfuller
    from statsmodels.tsa.stattools import adfuller

    # Calculate rolling statistics
    roll_mean = TS.rolling(window=5, center=False).mean()
    roll_std = TS.rolling(window=5, center=False).std()

    # Perform the Dickey Fuller test
    dfctest = adfuller(TS)

    # Plot rolling statistics:
    fig = plt.figure(figsize=(12,6))
    orig = plt.plot(TS, color='blue',label='Original')
    mean = plt.plot(roll_mean, color='red', label='Rolling Mean')
    std = plt.plot(roll_std, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show(block=False)

    # Print Dickey-Fuller test results
```

```

print('Results of Dickey-Fuller Test: \n')

dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value',
                                         '#Lags Used', 'Number of Observations Used'])

for key, value in dfctest[4].items():
    dfoutput['Critical Value (%s)'%key] = value
print(dfoutput)

return None

import statsmodels.api as sm
# Detrending the specified dataframes
# List of dataframe names to detrend
dataframe_names = [10011, 10021, 31561, 81611, 34102, 11975, 81615, 89413, 96141]

# Loop over the dataframe names
for dataframe_name in dataframe_names:
    # Get the dataframe for the dataframe name
    dataframe = region_dataframes[dataframe_name]

    # Calculate the difference between each observation and its value 12 months ago
    data_diff = dataframe.diff(periods=1)

    # Drop the missing values
    data_diff.dropna(inplace=True)

    ## Plot the detrended data
    fig = plt.figure(figsize=(11, 6))
    plt.plot(data_diff, color='blue', label='Detrended Values')
    plt.title(f"Detrended Values for Dataframe {dataframe_name}")
    plt.xlabel('Time')
    plt.ylabel('Detrended Values')
    plt.legend()
    plt.show()

    # Perform the stationarity check
    print(f"Stationarity Check for Dataframe {dataframe_name}:")
    stationarity_check(data_diff)
    #adf_result = sm.tsa.stattools.adfuller(data_diff)
    #p_value = adf_result[1]
    #print('P_value:',p_value)
    print()

```

Show hidden output

▼ Performing Detrending of our dataset and storing in a new dictionary

```

# Import adfuller
from statsmodels.tsa.stattools import adfuller

## Create a dictionary to store the detrended dataframes
detrended_dataframes = {}

dataframe_names = [10011, 10021, 31561, 81611, 34102, 11975, 81615, 89413, 96141]

# Loop over the dataframe names
for dataframe_name in dataframe_names:
    # Get the dataframe for the dataframe name
    dataframe = region_dataframes[dataframe_name]

    # Calculate the difference between each observation and its value 12 months ago
    data_diff = dataframe.diff(periods=1)

    # Drop the missing values
    data_diff.dropna(inplace=True)

    ## Plot the detrended data
    fig = plt.figure(figsize=(11, 6))
    plt.plot(data_diff, color='blue', label='Detrended Values')
    plt.title(f"Detrended Values for Dataframe {dataframe_name}")
    plt.xlabel('Time')
    plt.ylabel('Detrended Values')

```

```
plt.legend()
plt.show()

# Perform the stationarity check
def stationarity_check(TS):
    # Calculate rolling statistics
    roll_mean = TS.rolling(window=5, center=False).mean()
    roll_std = TS.rolling(window=5, center=False).std()

    # Perform the Dickey Fuller test
    dfctest = adfuller(TS)



    # Plot rolling statistics:
    fig = plt.figure(figsize=(12, 6))
    orig = plt.plot(TS, color='blue', label='Original')
    mean = plt.plot(roll_mean, color='red', label='Rolling Mean')
    std = plt.plot(roll_std, color='black', label='Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show(block=False)

    # Print Dickey-Fuller test results
    print('Results of Dickey-Fuller Test: \n')
    dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value', '#Lags Used', 'Number of Observations Used'])
    for key, value in dfctest[4].items():
        dfoutput['Critical Value (%s)' % key] = value
    print(dfoutput)

    return None


# Call the stationarity_check function
stationarity_check(data_diff)



# Store the detrended dataframe in the detrended_dataframes dictionary
detrended_dataframes[dataframe_name] = data_diff
```



Results of Dickey-Fuller Test:

Test Statistic	2.625870
p-value	0.999079
#Lags Used	10.000000
Number of Observations Used	25.000000
Critical Value (1%)	-3.723863
Critical Value (5%)	-2.986489
Critical Value (10%)	-2.632800
dtype: float64	








Results of Dickey-Fuller Test:

Test Statistic	-1.093233
p-value	0.717763
#Lags Used	7.000000
Number of Observations Used	28.000000
Critical Value (1%)	-3.688926
Critical Value (5%)	-2.971989
Critical Value (10%)	-2.625296



dtype: float64






Results of Dickey-Fuller Test:

Test Statistic	-3.532639
p-value	0.007187
#Lags Used	1.000000
Number of Observations Used	34.000000
Critical Value (1%)	-3.639224
Critical Value (5%)	-2.951230
Critical Value (10%)	-2.614447
dtype:	float64







Results of Dickey-Fuller Test:


Test Statistic	0.235991
p-value	0.974209
#Lags Used	8.000000
Number of Observations Used	27.000000
Critical Value (1%)	-3.699608
Critical Value (5%)	-2.976430
Critical Value (10%)	-2.627601
dtype:	float64

Detrended Values for Dataframe 34102



Rolling Mean & Standard Deviation







Results of Dickey-Fuller Test:

Test Statistic	-2.949263
p-value	0.039912
#Lags Used	3.000000
Number of Observations Used	32.000000
Critical Value (1%)	-3.653520
Critical Value (5%)	-2.957219
Critical Value (10%)	-2.617588
dtype:	float64

Detrended Values for Dataframe 11975



Rolling Mean & Standard Deviation





Results of Dickey-Fuller Test:


```

Test Statistic      -4.431406
p-value            0.000261
#Lags Used         1.000000
Number of Observations Used 34.000000
Critical Value (1%) -3.639224
Critical Value (5%) -2.951230
Critical Value (10%) -2.614447
dtype: float64

```






Results of Dickey-Fuller Test:

```

Test Statistic      0.605864
p-value            0.987759
#Lags Used         5.000000
Number of Observations Used 30.000000
Critical Value (1%) -3.669920
Critical Value (5%) -2.964071
Critical Value (10%) -2.621171
dtype: float64


```





Results of Dickey-Fuller Test:

Test Statistic	-3.585082
p-value	0.006053
#Lags Used	5.000000
Number of Observations Used	30.000000
Critical Value (1%)	-3.669920
Critical Value (5%)	-2.964071
Critical Value (10%)	-2.621171
dtype:	float64





```
#Displaying the detrended dataframe keys
detrended_dataframes.keys()
```

Show hidden output

Printing Mean & Standard Deviation

```
# Iterating over the items in the detrended_dataframes dictionary
for dataframe_name, detrended_dataframe in detrended_dataframes.items():
    print(f"Dataframe Name: {dataframe_name}")
    print(detrended_dataframe.head()) # Print the head of the detrended dataframe
print()
```

```
Dataframe Name: 10011
      value
```

```
time
```

```
2015-05-01   -3000.0
2015-06-01   105400.0
2015-07-01   173200.0
2015-08-01   283100.0
2015-09-01   211500.0
```

```
Dataframe Name: 10021
      value
```

```
time
```

```
2015-05-01   15500.0
2015-06-01   489700.0
2015-07-01   626000.0
2015-08-01   190600.0
2015-09-01  -112700.0
```

```
Dataframe Name: 31561
      value
```

```
time
```

```
2015-05-01   28900.0
2015-06-01   28400.0
2015-07-01   19300.0
2015-08-01    700.0
2015-09-01   11500.0
```

```
Dataframe Name: 81611
      value
```

```
time
```

```
2015-05-01   45000.0
2015-06-01   12800.0
2015-07-01   14600.0
2015-08-01   29900.0
2015-09-01   14800.0
```

```
Dataframe Name: 34102
      value
```

```
time
```

```
2015-05-01   20700.0
2015-06-01   19900.0
2015-07-01    7600.0
2015-08-01   11900.0
2015-09-01   17800.0
```

```
Dataframe Name: 11975
      value
```

```
time
```

```
2015-05-01   27900.0
2015-06-01   23200.0
2015-07-01   19300.0
2015-08-01   23900.0
2015-09-01   37300.0
```

```
Dataframe Name: 81615
      value
```

```
time
```

```
2015-05-01   26400.0
2015-06-01   20900.0
2015-07-01   23400.0
2015-08-01    6400.0
2015-09-01  -32100.0
```

```
Dataframe Name: 89413
      value
```

```
# Accessing the separate dataframes for each region name
for region_name, region_df in region_dataframes.items():
    print(f'Dataframe name: {region_name}')
    print(region_df.head())
    print()
```

```
Dataframe name: 10011
      time  value
```

```
2015-04-01  10572500.0
2015-05-01  10569500.0
2015-06-01  10674900.0
2015-07-01  10848100.0
2015-08-01  11131200.0
```

```
Dataframe name: 10014
      time  value
```

```
2015-04-01  9938600.0
2015-05-01  9827500.0
2015-06-01  9571200.0
2015-07-01  9278700.0
2015-08-01  9252000.0
```

```
Dataframe name: 10021
      time  value
```

```
2015-04-01  16644000.0
2015-05-01  16659500.0
2015-06-01  17149200.0
2015-07-01  17775200.0
2015-08-01  17965800.0
```

```
Dataframe name: 11975
      time  value
```

```
2015-04-01  2343300.0
2015-05-01  2371200.0
2015-06-01  2394400.0
2015-07-01  2413700.0
2015-08-01  2437600.0
```

```
Dataframe name: 31561
      time  value
```

```
2015-04-01  2453900.0
2015-05-01  2482800.0
2015-06-01  2511200.0
2015-07-01  2530500.0
2015-08-01  2531200.0
```

```
Dataframe name: 34102
      time  value
```

```
2015-04-01  2481500.0
2015-05-01  2502200.0
2015-06-01  2522100.0
2015-07-01  2529700.0
2015-08-01  2541600.0
```

```
Dataframe name: 81611
      time  value
```

```
2015-04-01  3956500.0
```

▼ Performing Deseasonalizing of our dataset and storing in a new dictionary

```
#Deseasonalizing the dataset
```

```
deseasonalized_data = {}
```

```
# Iterate over each key (region) in the region_dataframes dictionary
```

```
for region_name, data_diff in detrended_dataframes.items():
```

```
    # Time Series Decomposition
```

```
    result_mul = seasonal_decompose(data_diff['value'], model='additive', extrapolate_trend='freq', period=int(len(data_diff) / 2))
```

```
    # Deseasonalize
```

```
    deseasonalized = data_diff['value'].values / result_mul.seasonal
```

```
    # Plot
```

```
plt.figure()
plt.plot(deseasonalized)
plt.title(f'Housing Values Deseasonalized - {region_name}', fontsize=14)
plt.xlabel('Years')
plt.ylabel('Value')
plt.xticks(rotation = 45)
plt.show()
```

```
# Store the deseasonalized values in the dictionary
deseasonalized_data[region_name] = deseasonalized
```

Show hidden output

```
#Printing the deseasonalized data keys
deseasonalized_data.keys()
```

```
dict_keys([10011, 10021, 31561, 81611, 34102, 11975, 81615, 89413, 96141])
```

Performing Seasonal decomposition on the dictionary

```
# Iterate over each dataframe in the region_dataframes dictionary
for region_name, region_df in deseasonalized_data.items():
    # Perform seasonal decomposition
    decomposition = seasonal_decompose(region_df)
```

```
# Gather the trend, seasonality, and residuals
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid
```


```
# Plot the gathered statistics
plt.figure(figsize=(12, 8))
plt.subplot(411)
plt.plot(region_df, label='Original', color='blue')
plt.legend(loc='best')
plt.title(f"Original - {region_name}")
```


```
plt.subplot(412)
plt.plot(trend, label='Trend', color='blue')
plt.legend(loc='best')
plt.title(f"Trend - {region_name}")
```


```
plt.subplot(413)
plt.plot(seasonal, label='Seasonality', color='blue')
plt.legend(loc='best')
plt.title(f"Seasonality - {region_name}")
```


```
plt.subplot(414)
plt.plot(residual, label='Residuals', color='blue')
plt.legend(loc='best')
plt.title(f"Residuals - {region_name}")
```


```
plt.tight_layout()
plt.show()
```











Checking for stationarity of residuals using Seasonal Decompose

```
# Iterate over each dataframe in the region_dataframes dictionary
for region_name, region_df in deseasonalized_data.items():
    # Perform seasonal decomposition
    decomposition = seasonal_decompose(region_df)

    # Gather the residuals
    residual = decomposition.resid

    # Drop missing values from residuals
    ts_log_decompose = residual
    ts_log_decompose.dropna(inplace=True)

    # Perform the stationarity check
    print(f"Stationarity Check for {region_name}:")
    stationarity_check(ts_log_decompose)
    print()
```

Show hidden output

```
#Creating separate dictionaries for the individual zipnames
ts_10011 = deseasonalized_data[10011]
ts_10014 = region_dataframes[10014]
ts_10021 = deseasonalized_data[10021]
ts_81611 = deseasonalized_data[81611]
ts_31561 = deseasonalized_data[31561]
ts_34102 = deseasonalized_data[34102]
ts_81611 = deseasonalized_data[81611]
ts_81615 = deseasonalized_data[81615]
ts_89413 = deseasonalized_data[89413]
ts_96141 = deseasonalized_data[96141]
```

PLotting for the seasonal decomposition of the stationary Region Name 10014

```
# Perform seasonal decomposition
decomposition = seasonal_decompose(ts_10014)

# Gather the residuals
```